Analysis of Large-Scale Co-authorship Networks

Navigating Data Skew, Distributed Memory Constraints, and Matrix Approximations

ARSH SAJID VORA

Algorithms for Massive Datasets

March 2026

Declaration

I/We declare that this material, which I/We now submit for assessment, is entirely my/our own work and has not been taken from the work of others, save and to the extent that such work has been cited and acknowledged within the text of my/our work, and including any code produced using generative AI systems. I/We understand that plagiarism, collusion, and copying are grave and serious offences in the university and accept the penalties that would be imposed should I engage in plagiarism, collusion or copying. This assignment, or any part of it, has not been previously submitted by me/us or any other person for assessment on this or any other course of study.

Note regarding external platforms: The raw dataset used in this project was downloaded via the Kaggle API. However, the data engineering pipelines, the transition from unipartite to bipartite graph structures, and the matrix approximation algorithms discussed herein were developed independently based on course methodologies. To establish a high-speed computational baseline, we implemented the T^2-Approx algorithm, a recently proposed method by Wu et al. (2024) for efficient PageRank approximation on large graphs. While our primary focus remained on the distributed Spark implementations taught in the course, adapting this novel matrix factorization approach allowed us to contrast traditional Map-Reduce limitations with modern linear algebra optimizations.

Contents

1 Introduction **4**

2 The Chosen Dataset and Structural Considerations **4**

2.1 Data Ingestion and Scale 4

2.2 Considered Data Fields and Entity Resolution 4

3 Graph Architectures and Pre-processing Techniques **5**

3.1 The Unipartite Nightmare and the $O(N^2)$ Explosion 5

3.2 Pre-Processing 1: The Degree-Filter Compromise 5

3.3 Pre-Processing 2: The Bipartite Paradigm Shift 6

4 The Considered Algorithms and Their Implementations **7**

4.1 The Mathematics of Distributed PageRank 7

4.2 Implementation 1: Filtered Spark Map-Reduce 8

4.3 Implementation 2: Bipartite Spark Graph Logic 8

4.4 Implementation 3: T^2-Approx Matrix Approximation 9

5 Scalability with Data Size **9**

6 Description of Experiments and Results **10**

6.1 Experiment 1: The Matrix Baseline (T^2-Approx) 10

6.2 Experiment 2: The Filtered Spark Run 10

6.3 Experiment 3: The Bipartite Spark Run (Ground Truth) 10

7 Discussion and Comments on Experimental Results **11**

7.1 The "Entity Resolution" Problem 11

7.2 The Dominance of Big Science 11

7.3 Score Normalization Differences 11

8 Advanced Algorithmic and Execution Analysis **11**

8.1 Spark Execution Mechanics: Lineage and Shuffling 12

8.1.1 Wide Dependencies and The Shuffle Bottleneck 12

8.2 Mathematical Deep Dive: T^2-Approx Matrix Construction 12

8.2.1 Degree-Weighted Sampling 12

8.2.2 Constructing X and Y 13

9 Hardware Constraints and Configuration Tuning **13**

10 Final Conclusion and Project Reflections **14**

1 Introduction

In the era of modern academia, research has shifted from solitary endeavors to massive, interconnected collaborations. Understanding the topology of these collaborations allows us to identify the "hubs" of scientific progress—researchers and consortiums that act as foundational pillars for their respective fields. The computational challenge arises when we attempt to map this "Republic of Science" at scale.

The objective of this project was to implement the PageRank algorithm—originally designed to evaluate the hyperlink structure of the World Wide Web—on a massive dataset of academic authors. By treating authors as nodes and co-authorships as directed links, we aimed to calculate the stationary distribution of a "random surfer" navigating the academic network.

However, as the dataset scaled to 1,965,563 unique authors and 2,975,294 papers, naive graph representations led to catastrophic memory failures. This project evolved into an exercise in Big Data engineering, forcing a critical examination of how data structures interact with distributed hardware limits. This report details the journey through three distinct methodologies:

1. A degree-filtered Apache Spark RDD approach (Unipartite Graph).
2. An optimized, unfiltered Apache Spark approach (Bipartite Graph).
3. A high-speed Matrix Approximation (T^2-Approx) using SciPy.

2 The Chosen Dataset and Structural Considerations

2.1 Data Ingestion and Scale

The dataset utilized is the **ArXiv Metadata OAI Snapshot**, hosted by Cornell University on Kaggle. The dataset is approximately 1.58 GB in size and is structured as a series of JSON objects, where each line represents a single scholarly article.

To simulate a realistic cluster environment on a single-node Google Colab instance, we initialized Apache Spark with strict memory allocations. We assigned 8GB to the driver and 2GB to the executor to maximize the 12GB hardware limit.

2.2 Considered Data Fields and Entity Resolution

Each JSON record contains extensive metadata, including titles, abstracts, and categories. To prevent memory exhaustion during the Resilient Distributed Dataset (RDD) creation phase, we discarded the text-heavy fields and strictly isolated two attributes:

- `id`: The unique alphanumeric string representing the paper.

- `authors_parsed`: A structured list containing the separated first and last names of every contributing author.

The Entity Resolution Challenge: Relying on `authors_parsed` introduced a classic data science challenge: Name Disambiguation. Without unique researcher IDs (such as ORCID), our system treated identical string representations as the same node. For example, the string `"Liu, Yang"` appears thousands of times across different papers. The system merges all these distinct researchers into a single "Super-Node," artificially inflating their PageRank score. This phenomenon was heavily observed in our filtered experiments.

3 Graph Architectures and Pre-processing Techniques

The organization of the graph data proved to be the most critical bottleneck of this project. A network of this size requires careful structuring to avoid combinatorial explosions in the adjacency list.

3.1 The Unipartite Nightmare and the $O(N^2)$ Explosion

Initially, we attempted to map the data as a standard unipartite directed graph. The logic is straightforward: if Author A, Author B, and Author C write a paper together, we create bidirectional edges between all of them.

Mathematically, a paper with n authors generates a fully connected clique. The number of directed edges E generated by a single paper is:

$$E = n(n-1) \tag{1}$$

For a standard computer science paper with 4 authors, $E = 12$ edges. However, High-Energy Physics (HEP) papers routinely feature massive collaborations. The ATLAS and CMS collaborations can have over 3,000 authors on a single paper. For $n = 3,000$:

$$E = 3000(2999) = 8,997,000 \text{ edges} \tag{2}$$

When Spark attempted to `flatMap` these edges into memory, the JVM Garbage Collector immediately failed due to severe "Data Skew." The hardware simply could not shuffle millions of edges generated by a single JSON line.

3.2 Pre-Processing 1: The Degree-Filter Compromise

To force the unipartite graph to execute without crashing, we implemented a strict "Survival Filter." We parsed the JSON and explicitly discarded any paper with more than 50

authors.

Listing 1: Unipartite Edge Mapping with Degree Filter

```python
def parse_coauthors(line):
    try:
        record = json.loads(line)
        authors = record.get('authors_parsed', [])

        if len(authors) > 50 or len(authors) < 2:
            return []

        author_names = [f"{a[0].strip()}, {a[1].strip()}" for a in authors]
        edges = []

        for i in range(len(author_names)):
            for j in range(len(author_names)):
                if i != j:
                    edges.append((author_names[i], author_names[j]))
        return edges
    except:
        return []
```

While this code successfully prevented memory crashes and allowed the Spark PageRank to finish in approximately 90 minutes, it introduced a severe academic bias. By explicitly capping the degree at 50, we artificially erased the most collaborative fields of science (like High-Energy Physics) from our network entirely.

3.3 Pre-Processing 2: The Bipartite Paradigm Shift

To process 100% of the data accurately, we had to rethink the graph's topology. Instead of mapping Author → Author, we reorganized the data into a **Bipartite Graph**. In this structure, both *Authors* and *Papers* act as individual nodes.

Edges strictly flow between an author and the paper they wrote. We generate two directed edges per author:

- Author → Paper
- Paper → Author

Listing 2: Bipartite Edge Mapping (No Filters)

```python
def parse_bipartite_edges(line):
    try:
        record = json.loads(line)
        paper_id = f"P:{record.get('id')}"
        authors = record.get('authors_parsed', [])
        edges = []

        for a in authors:
            author_name = f"A:{a[0].strip()}, {a[1].strip()}"
            edges.append((author_name, paper_id))
            edges.append((paper_id, author_name))
        return edges
    except:
        return []
```

The Mathematical Breakthrough: By introducing the paper as an intermediary node, the number of edges generated scales linearly, $O(n)$. For the same 3,000-author ATLAS paper, the bipartite model generates exactly $2n = 6,000$ edges. This massive reduction in computational complexity allowed us to ingest all 1.96 million authors and 2.97 million papers without a single memory exception.

4 The Considered Algorithms and Their Implementations

4.1 The Mathematics of Distributed PageRank

The core logic of our graph analysis relies on the PageRank algorithm. Originally conceptualized to rank World Wide Web pages, it models a "random surfer" navigating a network. In our context, the surfer randomly clicks through co-authorship links.

The rank of a node u at iteration $k+1$ is calculated as:

$$PR(u) = \frac{1-\alpha}{N} + \alpha \sum_{v \in B(u)} \frac{PR(v)}{L(v)} \tag{3}$$

Where:

- α is the damping factor. Following standard academic practice, we set $\alpha = 0.85$. This implies an 85% chance the surfer follows a collaboration link, and a 15% chance they randomly teleport to any other author in the database.
- N is the total number of nodes in the graph.

- $B(u)$ is the set of nodes linking to u.
- $L(v)$ is the out-degree of node v (the number of co-authors they have).

4.2 Implementation 1: Filtered Spark Map-Reduce

Our first full implementation utilized PySpark's Resilient Distributed Datasets (RDDs). Because PageRank is an iterative algorithm, it is notoriously heavy on disk I/O if not managed correctly.

We initialized the ranks of all authors to 1.0. To prevent Spark from re-reading the 1.58GB JSON dataset on every single iteration, we utilized the professor's recommended optimizations: `partitionBy()` and `cache()`.

Listing 3: Spark RDD Optimization and Caching

```

links = edges_rdd.groupByKey().partitionBy(8).cache()

ranks = links.mapValues(lambda v: 1.0)
```

Setting `partitionBy(8)` explicitly mapped the RDD partitions to the 8 available CPU threads in our Colab environment. During the 10 iterations, we used a `flatMap` to distribute the rank fragments to neighbors, followed by a `reduceByKey` to sum the incoming fragments.

4.3 Implementation 2: Bipartite Spark Graph Logic

Processing the Bipartite Graph required a logical shift. Because the edges flow Author → Paper → Author, a single iteration of PageRank only moves the "rank" from an Author to a Paper. It takes a second iteration to move the rank from the Paper back to the co-authors.

Therefore, running 10 iterations on a Bipartite Graph effectively simulates 5 full "surfer hops" between authors.

Listing 4: Bipartite Spark PageRank Loop

```
for iteration in range(iterations):
    contributions = links.join(ranks).flatMap(
        compute_contributions)

    from operator import add
    ranks = contributions.reduceByKey(add) \
            .mapValues(lambda rank: (1 - 0.85) + 0.85 * rank)
```

Because the bipartite graph maintained $O(N)$ linear edge scaling, this implementation successfully processed 100% of the data.

4.4 Implementation 3: T^2-Approx Matrix Approximation

To establish a high-speed computational baseline, we implemented the T^2-Approx algorithm using SciPy sparse matrices. Spark is robust, but the JVM overhead of serializing Python objects slows it down significantly on a single node.

Instead of iterating over the massive 23-million edge matrix, we used degree-weighted sampling to select a small subset of columns, specifically $c = 3000$.

We constructed two highly compressed matrices, X and Y. Power iteration was then performed on the shrunken matrix Z:

$$Z = Y \cdot X \tag{4}$$

By projecting the graph into this lower-dimensional space, the vector converged rapidly. We set a dynamic stopping threshold of $\tau = \frac{1.0}{100 \times N_{authors}}$. The loop monitored the L^2-norm difference between iterations:

Listing 5: Power Iteration Convergence Check

```
if np.linalg.norm(pi_next - pi_k, ord=2) < tau:
    break
```

This matrix approximation bypassed the distributed shuffle overhead entirely, allowing for extremely rapid execution.

5 Scalability with Data Size

The scalability of our proposed solutions varied drastically based on the underlying architecture and the memory constraints of our 12GB RAM environment.

- **Unipartite Spark:** Highly unscalable. The $O(N^2)$ complexity caused massive memory pressure. When the shuffle size exceeded the 8GB driver limit, Spark began spilling temporary data to the local disk, causing execution times to stretch exponentially.
- **Bipartite Spark:** Highly scalable. The $O(N)$ edge growth meant that even as the dataset scales to tens of millions of papers, the memory footprint increases linearly. Spark effectively managed the RDD lineage without heavy disk-spilling.
- **Matrix Approximation (T^2-Approx):** Extremely fast, but fragile. Constructing the initial `scipy.sparse` matrix required the entire adjacency list to fit into

RAM simultaneously. It lacks Spark's robust "spill-to-disk" safety nets, meaning it would crash immediately if applied to a dataset 10x larger.

6 Description of Experiments and Results

We executed three isolated experiments to compare the methodologies.

6.1 Experiment 1: The Matrix Baseline (T^2-Approx)

We ran the matrix approximation on the unfiltered dataset. The power iteration converged in just 10 iterations. The total pipeline execution time, including data ingestion, was a staggering 84.69 seconds. The top result was the **ATLAS Collaboration**.

6.2 Experiment 2: The Filtered Spark Run

We executed the standard Spark Map-Reduce PageRank, filtering out any paper with more than 50 authors. The computation completed in approximately 90 minutes. The top results were dominated by individuals:

1. **Liu, Yang** (Score: 285.2601)
2. **Taniguchi, Takashi** (Score: 243.9999)
3. **Watanabe, Kenji** (Score: 236.8832)

6.3 Experiment 3: The Bipartite Spark Run (Ground Truth)

This run utilized the Bipartite model on 100% of the data without degree filters. Despite processing significantly more authors and papers than Experiment 2, it completed much faster, running in approximately 50 minutes. The true "hubs" of the network emerged:

1. **CMS Collaboration** (Score: 532.1447)
2. **ATLAS Collaboration** (Score: 498.9011)
3. **Liu, Yang** (Score: 357.1985)

7 Discussion and Comments on Experimental Results

7.1 The "Entity Resolution" Problem

In Experiment 2 (Filtered Spark), the top author identified was **Liu, Yang**. While mathematically correct based on our parsed graph, this reveals a major data science caveat: Name Disambiguation. Because we lacked unique researcher IDs, Spark treated every occurrence of "Liu, Yang" as a single node. The algorithm essentially merged hundreds of distinct researchers into one massive "Super-Node," artificially inflating its rank. The prominence of Asian surnames in the top 20 reflects this high-entropy name collision, compounded by the heavy research output in fields like Computer Science which typically have smaller co-author counts (thus surviving our 50-author filter).

7.2 The Dominance of Big Science

In Experiment 3 (Bipartite Spark), the High-Energy Physics community reclaimed the top spots. **CMS** and **ATLAS** are massive consortiums. In a PageRank context, they possess thousands of incoming links from their internal members, making them inescapable "gravitational centers" of the ArXiv network. Experiment 3 proves that mitigating Data Skew through architectural restructuring (Bipartite) rather than data deletion (Degree Filtering) yields the only scientifically accurate representation of the network.

7.3 Score Normalization Differences

A critical observation must be made regarding the mathematical scales of our results. In the T^2-Approx experiment, the PageRank vector was explicitly normalized so the sum of all ranks equaled 1.0, resulting in tiny fractional scores like 0.00009. Conversely, the Spark implementations initialized ranks at 1.0, meaning the total sum of the network equaled N (the total number of nodes). As influence flowed, the top hubs accumulated hundreds of points (e.g., 532.14). Both approaches correctly identify the stationary distribution, simply represented on different probabilistic scales.

8 Advanced Algorithmic and Execution Analysis

To fully demonstrate mastery over the processing of massive datasets, we must look beyond the outputs and analyze the execution mechanics of both the Apache Spark DAG and the Linear Algebra approximation.

8.1 Spark Execution Mechanics: Lineage and Shuffling

Apache Spark's core abstraction is the Resilient Distributed Dataset (RDD). In our PageRank implementations, understanding the difference between *Narrow* and *Wide* dependencies was the key to optimizing the 90-minute and 50-minute runtimes.

8.1.1 Wide Dependencies and The Shuffle Bottleneck

In our filtered unipartite graph, the `flatMap` operation that parses the JSON into edges is a narrow dependency—each partition of the input file can be processed independently. However, the PageRank algorithm requires grouping these edges by author to calculate out-degrees.

The operation `groupByKey()` induces a *Wide Dependency*. This triggers a "Shuffle," where Spark must reorganize data across the network so that all edges belonging to "Author A" end up on the same physical CPU core. In our Google Colab environment, standard shuffles of the 23-million edge dataset caused intense memory pressure.

To mitigate this, we explicitly controlled the partitioning:

Listing 6: Explicit Partitioning to Control Shuffles

```
# Matches Colab CPU threads to prevent over-partitioning
conf.set("spark.default.parallelism", "8")

# Forcing the graph topology into a known partitioned state
links = edges_rdd.groupByKey().partitionBy(8).cache()
```

By applying `.partitionBy(8)`, we forced Spark to hash-partition the nodes across the 8 available threads. Because we immediately called `.cache()`, this highly expensive shuffle was only performed *once*. During the 10 iterative loops, Spark only shuffled the much smaller 'ranks' RDD, joining it against the static, partitioned 'links' RDD in memory.

8.2 Mathematical Deep Dive: T^2-Approx Matrix Construction

While Spark handles scale through distribution, our T^2-Approx implementation handles scale through mathematical compression. The full transition matrix T for the ArXiv network has dimensions $N \times N$, where $N \approx 1.96 \times 10^6$. Storing this densely would require terabytes of RAM.

Instead of computing the exact 2-step transition matrix T^2, we constructed low-rank approximations X and Y such that $T^2 \approx YX$.

8.2.1 Degree-Weighted Sampling

To ensure the approximation accurately captures the most important hubs (like the ATLAS Collaboration), we did not sample columns uniformly. We used **Degree-Weighted**

Sampling. We calculated the degree of every paper d, and defined a probability distribution p_dist over the authors:

Listing 7: Degree-Weighted Probability Distribution

```
paper_degrees = A.T.dot(np.ones(n_authors, dtype=np.float32))
d = A.dot(paper_degrees)
d[d == 0] = 1.0
p_dist = d / np.sum(d)
```

We then randomly sampled $c = 3000$ authors based on this probability distribution to form a subset S.

$$S \sim \text{np.random.choice}(N, c, \text{p} = p_dist) \tag{5}$$

8.2.2 Constructing X and Y

Using the subset S, we extracted specific rows and columns from the adjacency matrix A to build our unnormalized intermediate matrices. We then applied the inverse degree matrix inv_d to normalize them:

$$Y_{unnormalized} = A_{S,:} \cdot A^T \tag{6}$$

$$X_{unnormalized} = A \cdot A^T_{S,:} \tag{7}$$

The final highly compressed transition matrix used for the Power Iteration was $Z = Y \cdot X$. By running the algorithm in this reduced $c \times c$ space, we achieved our 84.69-second pipeline time, bypassing the distributed overhead of PySpark entirely.

9 Hardware Constraints and Configuration Tuning

Running Big Data algorithms on a limited single-node instance requires deep tuning of the Java Virtual Machine (JVM) that underlies Apache Spark.

The Google Colab instance provided approximately 12GB of RAM. If we had instantiated Spark with its default parameters, the driver program would have been restricted to 1GB of heap space, causing an immediate `java.lang.OutOfMemoryError` during the JSON parsing phase.

We constructed a custom `SparkConf` object to push the hardware to its absolute limit:

Listing 8: Spark Hardware Tuning

```
conf = SparkConf() \
    .setAppName("Distributed_PageRank_Arxiv") \
    .setMaster("local[*]") \
    .set("spark.driver.memory", "8g") \
```

```
    .set("spark.executor.memory", "2g")
```

- `local[*]` instructs Spark to spin up a local cluster utilizing all available logical CPU cores.
- `spark.driver.memory = 8g` allocates the vast majority of the system RAM to the driver. This was mandatory because our `take(20)` actions pulled the aggregated results back to the driver node for sorting.
- `spark.executor.memory = 2g` provides just enough memory for the worker threads to process their partitions and perform local Map-Side Combinations before shuffling data across the network.

10 Final Conclusion and Project Reflections

This project serves as a comprehensive study in the processing of massive datasets. We successfully implemented and evaluated three distinct strategies for computing the PageRank of the 23-million edge ArXiv co-authorship network.

The journey from crashing Colab instances to achieving a stable, 50-minute execution highlighted the paramount importance of Graph Architecture. The traditional Unipartite Map-Reduce logic ($O(N^2)$), while theoretically sound, is practically unscalable for highly dense networks with massive power-law hubs (such as High-Energy Physics collaborations). Attempting to fix this via Degree Filtering artificially corrupted the academic reality of the graph.

The pivotal breakthrough was translating the data into a **Bipartite Graph** ($O(N)$), which effectively bounded the memory complexity and allowed Spark to process 100% of the data exactly. The results correctly and reliably identified the CMS and ATLAS Collaborations as the true hubs of modern science.

Conversely, the implementation of the T^2**-Approx** method proved that deep mathematical optimizations (Degree-Weighted Sampling and Matrix Factorization) can outperform distributed "brute force" computation by orders of magnitude (84 seconds vs. 50 minutes), albeit at the cost of exact precision and RAM stability.

Ultimately, mastering massive datasets requires fluency in both software engineering (Spark DAGs, memory tuning, JVM constraints) and advanced mathematics (Linear Algebra approximations, Random Surfer convergence). This project successfully bridged that gap, delivering a scalable, robust pipeline capable of mapping the most influential nodes in the Republic of Science.